

Building a Machine Learning Model for Heart Disease Prediction Dataset

Decision Tree

```
In [18]: # Import Libraries
import pandas as pd # Data manipulation
import numpy as np # Numerical operations
import matplotlib.pyplot as plt # Data visualization
import seaborn as sns # Data visualization
from sklearn.linear_model import LogisticRegression # Logistic Regression model
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score
from sklearn.model_selection import train_test_split # Train-test split
from sklearn.preprocessing import StandardScaler, LabelEncoder, MinMaxScaler # D
```

A. Data Understanding & Preprocessing:

```
In [19]: # Load and explore the dataset
df = pd.read_csv('heart_disease.csv')
# Display the first few rows of the dataset
print(df.head())
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	52	1	0	125	212	0	1	168	0	1.0	2	
1	53	1	0	140	203	1	0	155	1	3.1	0	
2	70	1	0	145	174	0	1	125	1	2.6	0	
3	61	1	0	148	203	0	1	161	0	0.0	2	
4	62	0	0	138	294	1	1	106	0	1.9	1	

	ca	thal	target
0	2	3	0
1	0	3	0
2	0	3	0
3	1	3	0
4	3	2	0

```
In [20]: # Check for missing values
df.isnull().sum()
```

```
Out[20]: age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```

```
In [21]: df.shape # Display the data types of each column
```

```
Out[21]: (1025, 14)
```

B. Exploratory Data Analysis (EDA)

```
In [23]: # Display information about the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1025 non-null   int64
 1   sex         1025 non-null   int64
 2   cp          1025 non-null   int64
 3   trestbps    1025 non-null   int64
 4   chol        1025 non-null   int64
 5   fbs         1025 non-null   int64
 6   restecg     1025 non-null   int64
 7   thalach     1025 non-null   int64
 8   exang       1025 non-null   int64
 9   oldpeak     1025 non-null   float64
10   slope       1025 non-null   int64
11   ca          1025 non-null   int64
12   thal        1025 non-null   int64
13   target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

```
In [24]: # Display summary statistics of the dataset
df.describe()
```

Out[24]:

	age	sex	cp	trestbps	chol	fbs	
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	10
mean	54.434146	0.695610	0.942439	131.611707	246.000000	0.149268	
std	9.072290	0.460373	1.029641	17.516718	51.59251	0.356527	
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	
25%	48.000000	0.000000	0.000000	120.000000	211.000000	0.000000	
50%	56.000000	1.000000	1.000000	130.000000	240.000000	0.000000	
75%	61.000000	1.000000	2.000000	140.000000	275.000000	0.000000	
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	

In [25]: *# removing duplicates*
`df = df.drop_duplicates()`

In [26]: `df.shape` *# Display the data types of each column*


Out[26]: (302, 14)

In [27]: `df.columns`

Out[27]: Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',
'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],
dtype='object')

In [28]: *# Visualize relationships between features*
`plt.figure(figsize=(15, 10))`
`sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt='.2f')`
`plt.title('Correlation Matrix of Heart Disease Features')`

Out[28]: Text(0.5, 1.0, 'Correlation Matrix of Heart Disease Features')



C. Feature Selection:

ONE-HOT ENCODING (Preparing Categorical Data)

Some of our columns have numbers that represent categories, not actual values. For example, 'cp' (chest pain type) has values 0, 1, 2, 3. These are just labels. A model might mistakenly think $3 > 1$, which is not true. We use One-Hot Encoding. This creates new columns for each category with a 1 or 0, telling the model if that category is present or not.

```
In [29]: # List the columns that are categorical
categorical_cols = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal']
```

```
In [30]: # Use pd.get_dummies to perform one-hot encoding
# `drop_first=True` is a good practice to avoid having redundant columns.
df_encoded = pd.get_dummies(df, columns=categorical_cols, drop_first=True)
```

```
In [31]: df_encoded.shape # Check the shape after encoding
```

```
Out[31]: (302, 23)
```

```
In [32]: df_encoded.head() # Display the first few rows of the encoded DataFrame
```

Out[32]:

	age	trestbps	chol	thalach	oldpeak	target	sex_1	cp_1	cp_2	cp_3	...	exang_1
0	52	125	212	168	1.0	0	True	False	False	False	...	False
1	53	140	203	155	3.1	0	True	False	False	False	...	True
2	70	145	174	125	2.6	0	True	False	False	False	...	True
3	61	148	203	161	0.0	0	True	False	False	False	...	False
4	62	138	294	106	1.9	0	False	False	False	False	...	False

5 rows × 23 columns

SPLIT DATA INTO TRAINING AND TESTING SETS

Training Set (80%): To teach/train our model. Testing Set (20%): To see how well our model performs on new, unseen data.

```
In [34]: # Define features and target variable
X = df_encoded.drop('target', axis=1)
y = df_encoded['target']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
```

D. Model Building & Evaluation

```
In [35]: # model call
model = LogisticRegression()
```

```
In [36]: # train the model
model.fit(X_train, y_train)
```

c:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\sklearn\linear_model_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

```
Out[36]: ▼ LogisticRegression ⓘ ?
LogisticRegression()
```

```
In [37]: # predict the values
y_pred = model.predict(X_test)
```

```
In [38]: y_pred # Evaluate the model
```

```
Out[38]: array([1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1,
          0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0,
          0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1])
```

EVALUATE THE MODEL'S PERFORMANCE

```
In [39]: # evaluate the model
print('Accuracy Score: ', accuracy_score(y_test, y_pred)) # Accuracy of the model
print('Recall Score: ', recall_score(y_test, y_pred)) # Recall of the model
print('Precision Score: ', precision_score(y_test, y_pred)) # Precision of the model
print('F1 Score: ', f1_score(y_test, y_pred)) # F1 Score of the model
print('Confusion Matrix: \n', confusion_matrix(y_test, y_pred)) # Confusion Matrix
print('Classification Report: \n', classification_report(y_test, y_pred)) # Classification Report
```

Accuracy Score: 0.7868852459016393

Recall Score: 0.8620689655172413

Precision Score: 0.7352941176470589

F1 Score: 0.7936507936507936

Confusion Matrix:

```
[[23  9]
```

```
[ 4 25]]
```


Classification Report:

	precision	recall	f1-score	support
0	0.85	0.72	0.78	32
1	0.74	0.86	0.79	29
accuracy			0.79	61
macro avg	0.79	0.79	0.79	61
weighted avg	0.80	0.79	0.79	61

confusion matrix

```
In [40]: # Plotting the confusion matrix
plt.figure(figsize=(8, 4))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Actual')
```

```
Out[40]: Text(70.7222222222221, 0.5, 'Actual')
```



Decision Tree Algorithm

Training the DecisionTreeClassifier

```
In [42]: from sklearn.tree import DecisionTreeClassifier

# create and train teh model with pred
model = DecisionTreeClassifier(criterion='entropy', random_state=42, max_depth=5)
model.fit(X_train, y_train)
```


```
Out[42]: DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=5, random_state=42)
```

```
In [43]: # predict the model
y_pred = model.predict(X_test)
# evaluate the model
print(confusion_matrix(y_test, y_pred)) # Confusion Matrix of the model
print(classification_report(y_test, y_pred)) # Classification Report of the model
```

```
[[18 14]
 [ 6 23]]
```

	precision	recall	f1-score	support
0	0.75	0.56	0.64	32
1	0.62	0.79	0.70	29
accuracy			0.67	61
macro avg	0.69	0.68	0.67	61
weighted avg	0.69	0.67	0.67	61

```
In [44]: # plot confusion matrix
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```



Convert into HTML, open the HTML file in a browser and print to PDF

```
In [46]: # Code to save this notebook as PDF
import subprocess
import os

def save_notebook_as_pdf():
    """
    Save the current notebook as PDF using nbconvert
    """
    try:
        # Get the current notebook path
        notebook_path = "Decision Tree.ipynb"

        # Method 1: Using nbconvert command line
        print("Attempting to save notebook as PDF...")

        # Install required packages if not available
        subprocess.run(["pip", "install", "nbconvert[webpdf]"], check=True)

        # Convert notebook to PDF
        result = subprocess.run([
            "jupyter", "nbconvert",
```



```

        "--to", "pdf",
        notebook_path
    ], capture_output=True, text=True)

    if result.returncode == 0:
        print(f"✅ Successfully saved as 'Decision Tree.pdf'")
    else:
        print(f"❌ Error: {result.stderr}")

    # Alternative method using HTML first
    print("Trying alternative method: Converting to HTML first...")

    # Convert to HTML
    html_result = subprocess.run([
        "jupyter", "nbconvert",
        "--to", "html",
        notebook_path
    ], capture_output=True, text=True)

    if html_result.returncode == 0:
        print("✅ Successfully saved as 'Decision Tree.html'")
        print("You can open the HTML file in a browser and print to PDF")
    else:
        print(f"❌ HTML conversion also failed: {html_result.stderr}")

except Exception as e:
    print(f"❌ Error occurred: {str(e)}")
    print("\n📋 Manual steps to save as PDF:")
    print("1. Go to File > Print Preview in VS Code")
    print("2. Or use: File > Export > Export as PDF")
    print("3. Or convert to HTML first: jupyter nbconvert --to html 'Decision Tree'")

# Run the function
save_notebook_as_pdf()

```

Attempting to save notebook as PDF...

✗ Error: [NbConvertApp] Converting notebook Decision Tree.ipynb to pdf

[NbConvertApp] ERROR | Error while converting 'Decision Tree.ipynb'

Traceback (most recent call last):

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\nbconvertapp.py", line 487, in export_single_notebook

output, resources = self.exporter.from_filename(

notebook_filename, resources=resources

)

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\exporters\templateexporter.py", line 390, in from_filename

return super().from_filename(filename, resources, **kw) # type:ignore[return-value]

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\exporters\exporter.py", line 201, in from_filename

return self.from_file(f, resources=resources, **kw)

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\exporters\templateexporter.py", line 396, in from_file

return super().from_file(file_stream, resources, **kw) # type:ignore[return-value]

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\exporters\exporter.py", line 220, in from_file

return self.from_notebook_node(

nbformat.read(file_stream, as_version=4), resources=resources, **kw

)

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\exporters\pdf.py", line 184, in from_notebook_node

latex, resources = super().from_notebook_node(nb, resources=resources, **kw)

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\exporters\latex.py", line 92, in from_notebook_node

return super().from_notebook_node(nb, resources, **kw)

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\exporters\templateexporter.py", line 429, in from_notebook_node

output = self.template.render(nb=nb_copy, resources=resources)

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\jinja2\environment.py", line 1295, in render

self.environment.handle_exception()

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\jinja2\environment.py", line 942, in handle_exception

raise rewrite_traceback_stack(source=source)

File "C:\Users\osama\miniconda3\envs\python_ml\share\jupyter\nbconvert\templates\latex\index.tex.j2", line 8, in top-level template code

((* extends cell_style *))

File "C:\Users\osama\miniconda3\envs\python_ml\share\jupyter\nbconvert\templates\latex\style_jupyter.tex.j2", line 176, in top-level template code

\prompt{(((prompt)))}{(((prompt_color)))}{(((execution_count)))}{(((extra_space)))}

```

^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
File "C:\Users\osama\miniconda3\envs\python_ml\share\jupyter\nbconvert\template
s\latex\base.tex.j2", line 7, in top-level template code
((%- extends 'document_contents.tex.j2' -*))
^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^

File "C:\Users\osama\miniconda3\envs\python_ml\share\jupyter\nbconvert\template
s\latex\document_contents.tex.j2", line 51, in top-level template code
((%- block figure scoped -*))
^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^

File "C:\Users\osama\miniconda3\envs\python_ml\share\jupyter\nbconvert\template
s\latex\display_priority.j2", line 5, in top-level template code
((%- extends 'null.j2' -*))
^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^

File "C:\Users\osama\miniconda3\envs\python_ml\share\jupyter\nbconvert\template
s\latex\null.j2", line 30, in top-level template code
((%- block body -*))
File "C:\Users\osama\miniconda3\envs\python_ml\share\jupyter\nbconvert\template
s\latex\base.tex.j2", line 241, in block 'body'
((( super() )))
~~~~~^~~~~~

File "C:\Users\osama\miniconda3\envs\python_ml\share\jupyter\nbconvert\template
s\latex\null.j2", line 32, in block 'body'
((%- block any_cell scoped -*))
^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^

File "C:\Users\osama\miniconda3\envs\python_ml\share\jupyter\nbconvert\template
s\latex\null.j2", line 85, in block 'any_cell'
((%- block markdowncell scoped-*)) ((%- endblock markdowncell -*))
^^^^^^^^^^

File "C:\Users\osama\miniconda3\envs\python_ml\share\jupyter\nbconvert\template
s\latex\document_contents.tex.j2", line 68, in block 'markdowncell'
((( cell.source | citation2latex | strip_files_prefix | convert_pandoc('markd
own+tex_math_double_backslash', 'json',extra_args=[]) | resolve_references | conv
ert_explicitly_relative_paths | convert_pandoc('json','latex'))))
^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\filt
ers\pandoc.py", line 36, in convert_pandoc
    return pandoc(source, from_format, to_format, extra_args=extra_args)
File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\util
s\pandoc.py", line 50, in pandoc
    check_pandoc_version()
    ~~~~~^~~~~~

File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\util
s\pandoc.py", line 98, in check_pandoc_version
    v = get_pandoc_version()
File "C:\Users\osama\miniconda3\envs\python_ml\Lib\site-packages\nbconvert\util
s\pandoc.py", line 75, in get_pandoc_version
    raise PandocMissing()
nbconvert.utils.pandoc.PandocMissing: Pandoc wasn't found.
Please check that pandoc is installed:
https://pandoc.org/installing.html

Trying alternative method: Converting to HTML first...
✔ Successfully saved as 'Decision Tree.html'
You can open the HTML file in a browser and print to PDF

```